



LUIZ AUGUSTO FIDALGO

Frankfurt, Germany - open to relocating | **Nationality:** Portuguese (EU Citizen) & Brazilian
luiz.fidalgo@hotmail.com | +49 15679 592599
luizfidalgo.com | linkedin.com/in/luiz-fidalgo | github.com/lafidalgo
DOB: 05.05.2000 | **Marital Status:** Married
Languages: English (C2), Portuguese (Native)

PROFILE

Machine Learning Engineer with 3 years of experience building production LLM, RAG, and AI agent systems across AWS and GCP, and delivering AI-automation workshops to enterprises.

EDUCATION & CERTIFICATIONS

Bachelor of Electrical Engineering (with Highest Honors) <i>University of Brasilia</i> Ranked 1st among all graduating Electrical Engineering students	Mar 2018 - Feb 2024 <i>Brasilia, Brazil</i>
International Exchange Program (ICT, Internet & Mechatronics Engineering) <i>University of Rome Tor Vergata</i> Completed master's-level coursework in Artificial Intelligence and Internet of Things (IoT)	Feb 2023 - Aug 2023 <i>Rome, Italy</i>
AWS Certified Machine Learning Engineer - Associate <i>Amazon Web Services</i>	Feb 2026 Verify

PROFESSIONAL EXPERIENCE

Machine Learning Engineer <i>ada Learning (corporate AI upskilling / EdTech)</i> - Developing an EU-hosted, no-code AI-agent platform across infrastructure (Terraform, Docker, CI/CD), back-end (Python/FastAPI), and frontend (SvelteKit) - Built multi-agent orchestration on CrewAI, provider-agnostic LLM access via LiteLLM (OpenAI, Anthropic, Bedrock), RAG (Qdrant), tool-calling via MCP, and agentic-workflow evaluation (Langfuse) - Delivered AI Agent workshops to 70+ professionals representing 10+ different companies	Feb 2025 - Present <i>Dusseldorf, Germany</i>
Machine Learning Engineer <i>Hotmart (Brazilian digital-products platform, 30M+ users)</i> - Developed async API using LlamaIndex, sustaining 300+ RPS of LLM, RAG, and agent traffic, enabling data scientists to run high-concurrency inference across the platform - Load-tested the API with Locust and monitored performance in New Relic to configure a cost-efficient horizontal-scaling strategy on smaller instances - Built CI/CD automation (Docker, Kubernetes, Terraform) that streamlined ML model deployment across AWS SageMaker, EKS, Lambda, and SQS	Aug 2024 - Jan 2025 <i>Brasilia, Brazil</i>
Machine Learning Engineer <i>snowfoxAI (AI consultancy)</i> - Developed document extraction pipeline for an insurance client using OCR, LLMs, and RAG, achieving 97% accuracy while reducing per-document processing time from 3 minutes to 30 seconds (deployed on AWS) - Built retail computer vision system detecting and classifying 150+ products per image (YOLOv8, SAM), with model training on Vertex AI and scalable deployment on Google Kubernetes Engine (GKE) - Partnered with business stakeholders to define AI roadmaps, shape project scope, and design AI and data architectures	Sep 2023 - Aug 2024 <i>Brasilia, Brazil</i>
IoT Development Intern <i>GrupoOrion (IoT solutions company)</i> - Developed 5 IoT sensor prototypes measuring environmental factors (sound, smell, weight, people counting) - Produced, tested, and deployed 200+ sensor units from prototype to production - Built software tools for remote sensor testing, monitoring, and configuration using Azure IoT Hub	Mar 2021 - Dec 2022 <i>Brasilia, Brazil</i>

TECHNOLOGIES & SKILLS

Languages: Python, C, C++, TypeScript, SQL
GenAI & LLM: CrewAI, LlamaIndex, LiteLLM, MCP, RAG, Qdrant, Langfuse
ML & Computer Vision: PyTorch, scikit-learn, YOLOv8, SAM, OCR
Cloud: AWS (SageMaker, EKS, Lambda, SQS, S3), GCP (Vertex AI, GKE)
DevOps & Observability: Docker, Kubernetes, Terraform, CI/CD (GitHub Actions), Locust, New Relic
Web & Data: FastAPI, Flask, Streamlit, SvelteKit, PostgreSQL, MongoDB, DynamoDB